

183.605

Machine Learning for Visual Computing

Assignment 1

Group 12:
Hanna Huber (0925230)
Lena Trautmann (1526567)
Elisabeth Wetzler ()

November 16, 2016

- Upload a zip-file with the required programs. You can choose the programming language.
- Add a PDF document with answers to all of the questions of the assignment (particularly all required plots) and description and discussion of results.

1 Assignment 1

1.1 Part 1: Binary classification and the perceptron

1.1.1 Reading data

Tasks:

- Read the data using functions of your programming language resp. simulation software.

First, we wrote a new .csv-file with our .py-script to delete all extra blanks. Afterwards we could read in the data using `dlmread()`.

- Plot the input vectors in \mathbb{R}^2 and visualize corresponding target values (e.g. by using color).

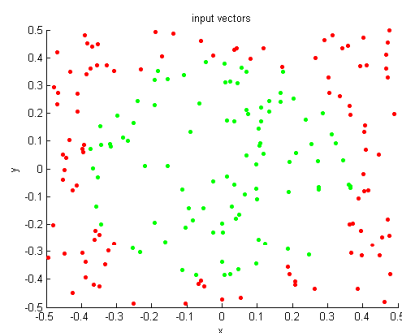


Figure 1: Plot of the input vectors with the target value visualized by colour.

Fig. 1 shows the input vectors.

- Use the feature transformation $(x_1, x_2) \rightarrow (x_1^2, x_2^2)$ and plot the data in the new feature space. The data should now be linearly separable.

Fig. 2 shows the transformed input vectors.

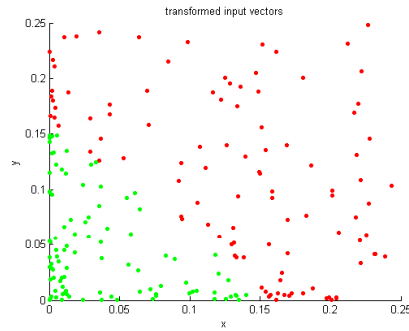


Figure 2: Plot of the transformed input vectors with the target value visualized by colour.

1.1.2 Perceptron training algorithm

The function

$$y = \text{perc}(w, X).$$

simulates a perceptron. The first argument is the weight vector w and the second argument is a matrix with input vectors in its columns X . The output y is a binary vector with class labels 1 or -1.

The function

$$w = \text{percTrain}(X, t, \text{maxIts}, \text{online}).$$

returns a weight vector w corresponding to the decision boundary separating the input vectors in X according to their target values t .

The argument `maxIts` determines an upper limit for iterations of the gradient based optimization procedure. If this upper limit is reached before a solution vector is found, the function returns the current w , otherwise it returns the solution weight vector. `online` is *true* if the *online*-version of the optimization procedure is to be used or *false* for the *batch*-version.

Tasks:

- Implement both functions. Use homogeneous coordinates and the corresponding augmented weight vector $w \in \mathbb{R}^3$.
- Plot the data and decision boundary in \mathbb{R}^2 , both in the original data space (see e.g. Figure 3) and in the feature space of basis functions, each at three different stages of the training: after the first iteration, after approximately half of the required iterations, and after convergence.
- Initialize $w = 0$. What is the influence of the learning rate?




Figure 3: Plot of the decision boundary in the original data space found by the perceptron (green curve) together with labelled data points.

Table 1: Weight vector derived for the LMS-rule and of closed form

	ϕ_0	ϕ_1	ϕ_2
wLMS	-0.1288	-16.6005	3.7258
wClosed			

1.2 Part 2: Linear basis function models for regression

1.2.1 Experimental setup

At the end of the setup we have:

- xtrain and ttrain: the training data
- phi: the transformation function
- xtrain_phi: the transformation of xtrain

1.2.2 Optimization: LMS-learning rule vs. closed form

Tasks:

- What is the resulting weight vector when using the LMS-rule?




Figure 4: Optimization using the LMS-rule

The values of the resulting weight vector are summarized in Tab. 1 and the corresponding curve plotted in Fig. 4.

- How can you determine the optimal \mathbf{w}^* in closed form? Compare \mathbf{w}^* with the outcome of the LMS-rule training.

The closed form is calculated by updating the weight vector, determined by all N points, in one step.

- What is the influence of γ ? Which value for γ represents a good tradeoff between number of iterations and convergence?

1.2.3 Model-complexity and model-selection

Determine \mathbf{w}^* in closed form for 2000 different training sets, in which only the t_i are varied according to $\mathcal{N}(\mu = y_i, \sigma = 16)$, while the x_i remain unchanged.

Tasks¹:

- Select a fixed x' , which is not an observation of the training set, but lies between two observations (e.g. $x' = 2$)
- Estimate the *mean squared error*

$$mse = \mathcal{E}(f(x') - f_{\mathbf{w}^*}(x'))^2, \quad (1)$$

i.e., the mean of the squared residuals of the models prediction $f_{\mathbf{w}^*}(x')$ from the true function value $f(x')$ for all $0 \leq d \leq 8$ ($d = 0$ corresponds to a constant function) using at least 2000 trials.

¹In all tasks \mathcal{E} refers to the expected value with respect to the random variable \mathbf{w}^* , i.e. $\mathcal{E} \equiv \mathcal{E}_{\mathbf{w}^*}$

- Estimate by the same way the quantities $bias^2 = (f(x') - \mathcal{E}f_{\mathbf{w}^*}(x'))^2$ and $var = (f_{\mathbf{w}^*}(x') - \mathcal{E}f_{\mathbf{w}^*}(x'))^2$.
- Plot mse , $bias^2$ and var against d together in one plot. What is the relation of the quantities?
- (optional) Generate the above plots only for $d = 8$, but minimize instead of $E(\mathbf{w})$ the regularized error function

$$E_{\lambda}(\mathbf{w}) = \sum_{i=1}^N (t_i - \mathbf{w}^T \Phi(x_i))^2 + \lambda \|\mathbf{w}\|^2, \quad (2)$$

i.e. $\mathbf{w}^* = \arg \min_{\mathbf{w}}^* E_{\lambda}(\mathbf{w})$. Plot the quantities against λ instead of d . Hint: The minimum can be obtained in closed form (see lecture slides).